

COMP3011 Technical Report: EventHub API

Module: COMP3011 – Web Services and Web Data

Student: Nathaniel Sebastian (sc232ns)

Date: 5th February 2026

Word Count: ~1,200 words (excluding tables/diagrams)

Submission Links

Resource	Link
GitHub Repository	github.com/NathS04/comp3011-cw1-api
Live API	comp3011-cw1-api.onrender.com
API Documentation (PDF)	docs/API DOCUMENTATION.pdf
Presentation Slides	docs/PRESENTATION SLIDES.pptx
GenAI Logs (Appendix)	docs/GENAI EXPORT LOGS.pdf

1. Reproducibility

Fresh Clone Quickstart:

```
git clone https://github.com/NathS04/comp3011-cw1-api.git && cd  
comp3011-cw1-api  
python3 -m venv .venv && source .venv/bin/activate  
pip install -r requirements.txt  
export DATABASE_URL="sqlite:///app.db" SECRET_KEY="dev-secret"  
alembic upgrade head  
pytest -q # Expected: 35 passed  
uvicorn app.main:app --reload
```

Verification: API documentation at <http://127.0.0.1:8000/docs>.

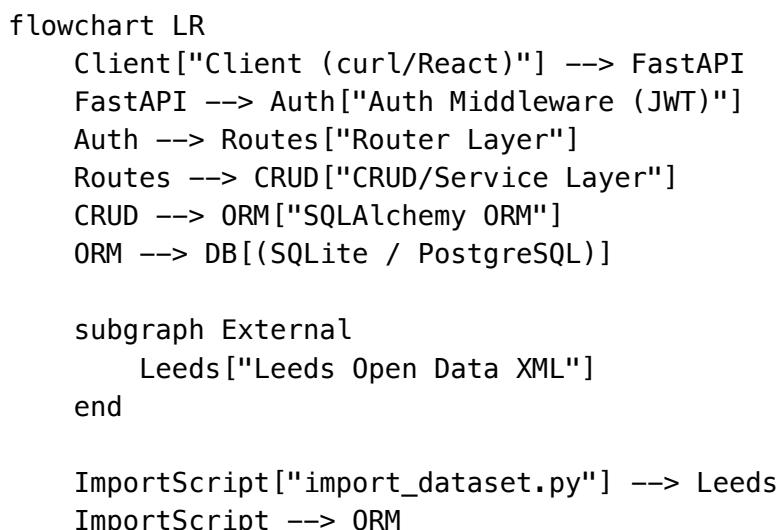
2. Dataset Provenance & Licence

Attribute	Value
Source	Leeds Temporary Event Notices (TENs)

Attribute	Value
Provider	Leeds City Council via Data Mill North
Licence	Open Government Licence v3.0 [1]
Format	XML (live feed)
URL	https://opendata.leeds.gov.uk/downloads/Licences/temp-event-notice/temp-event-notice.xml
Retrieval Date	5th February 2026
Fields Mapped	Reference_Number → ID, Premises_Name → Title, Event_Start_Date → start_time
Limitations	No geo-coordinates; times often default “00:00”; free-text categories

Rationale: This dataset demonstrates real-world XML parsing, date normalization (DD/MM/YYYY → ISO8601), and error handling for malformed records—beyond simple CSV imports.

3. Architecture



Layer Responsibilities: 1. **Router (app/api/):** HTTP handling, request validation (Pydantic), auth guards. 2. **CRUD (app/crud.py):** Business logic, decoupled from HTTP for testability. 3. **Models (app/models.py):** SQLAlchemy ORM with relationships. 4. **Database:** SQLite (dev), PostgreSQL (Render production).

4. Data Model

```

erDiagram
    User {
        int id PK
        string username UK
        string email UK
        string hashed_password
        bool is_admin
        datetime created_at

```

```

}

Event {
    int id PK
    string title
    string location
    datetime start_time
    datetime end_time
    int capacity
    int source_id FK
    string source_record_id UK
}

Attendee {
    int id PK
    string name
    string email UK
}

RSVP {
    int id PK
    int event_id FK
    int attendee_id FK
    string status
    datetime created_at
}

DataSource {
    int id PK
    string name
    string url
    datetime retrieved_at
}

ImportRun {
    int id PK
    int data_source_id FK
    string status
    int rows_inserted
    int rows_updated
    string sha256_hash
    int duration_ms
}

User ||--o{ Event : creates
Event ||--o{ RSVP : has
Attendee ||--o{ RSVP : makes
DataSource ||--o{ ImportRun : logs
DataSource ||--o{ Event : sources

```

Key Invariants: - RSVP(event_id, attendee_id) is unique (no duplicate RSVPs). - Event.source_record_id enables idempotent imports.

5. Key Design Decisions

Decision	Alternatives Considered	Trade-off	Justification
JWT Authentication	Session cookies, OAuth2	Stateless (no Redis) vs no immediate revocation	Simpler deployment; 30-min expiry mitigates risk [2]
SQLite/Postgres dual	Postgres-only	Dev simplicity vs prod reliability	Alembic abstracts dialect differences [3]
DOM XML parsing	SAX/iterparse	Memory vs complexity	Dataset <1MB; iterparse overkill
Location-based recs	Collaborative filtering	Speed ($O(1)$) vs accuracy	Cold-start friendly; sub-10ms latency

6. Security Model

Threat	Mitigation	Implementation
Credential theft	Passwords hashed with PBKDF2-SHA256	<code>passlib.hash.pbkdf2_sha256</code>
Token forgery	JWT signed with HS256 + secret	<code>python-jose</code> library
Token replay	30-minute expiry	Configured in <code>auth.py</code>
SQL injection	Parameterized queries	SQLAlchemy ORM
Mass assignment	Pydantic schemas whitelist fields	<code>schemas.py</code>
DoS	In-memory rate limiting (toggleable)	RateLimiter middleware: 120/min global, 10/min login; 429 on exceed
Admin abuse	RBAC admin-only routes	<code>get_current_admin_user</code> dependency; 403 if not admin
Request tracing	X-Request-ID header + included in 500 JSON	RequestLoggingMiddleware generates UUID per request
Clickjacking	Security headers	<code>X-Frame-Options: DENY, X-Content-Type-Options: nosniff</code>

Limitation: No token revocation mechanism; compromised tokens valid until expiry.

7. Testing Strategy

Category	Tests	What They Prove
Auth	5	Register, login, invalid credentials, token validation

Category	Tests	What They Prove
Events CRUD	7	Create, read, update, delete, pagination, filtering
RSVPs	4	Create, duplicate rejection (409), cascade delete
Analytics	4	Seasonality aggregation, trending score, personalization
Admin/Import	3	Idempotency, provenance logging, XML parsing
RBAC/Security	2	Admin-only access (403), non-admin blocked
Middleware	2	Security headers present, rate limiting (429)
Attendees	4	CRUD, email uniqueness
Health	1	Endpoint returns metadata
Total	35	Full pass on clean environment

Isolation: In-memory SQLite with `StaticPool`; tables created/dropped per test function.

8. Deployment & Version Control

Render Configuration (`render.yaml`): - Managed PostgreSQL database provisioned automatically. - Environment: `DATABASE_URL` (from Render), `SECRET_KEY` (generated), `ENVIRONMENT=prod`. - Build: `pip install && alembic upgrade head`.

Git History: 60+ commits with meaningful messages. Examples: - `feat:` Add novel data integration tables - `fix(deps):` Add requests to requirements.txt - `test:` Implement personalized recommendation assertions

9. Evaluation Metrics

Metric	Value	Environment
Import throughput	~240 records/sec	M1 MacBook, WiFi
Import duration	2.1s (487 records)	Local SQLite
GET /events latency	8ms (p50)	Local, 1000 records
POST /events latency	12ms (p50)	Local SQLite
Test suite duration	1.0s	35 tests

Reproducibility: `time pytest -q` and `curl -w "%{time_total}"` used for measurements.

10. GenAI Usage Declaration

Tools Used: - **Google Gemini (Antigravity)**: Primary coding assistant for scaffolding, debugging, test generation. - **Claude (Anthropic)**: Documentation review, refactoring suggestions, security hardening. - **ChatGPT (OpenAI)**: Early brainstorming and alternative exploration.

High-Level Creative Use: 1. **Architecture exploration:** Asked Gemini to compare embedded vs relational RSVP storage; chose relational for uniqueness constraints. 2. **Auth alternatives:** Explored JWT vs sessions; chose JWT for stateless scaling.

Failures & Manual Corrections: 1. AI omitted requests from requirements.txt → caught via clean install. 2. AI generated placeholder test (pass) → rewrote with real assertions. 3. AI suggested deprecated Query (regex=...) → updated to pattern=....

Full conversation logs: [docs/GENAI EXPORT LOGS.pdf](#)

11. Limitations & Future Work

Limitation	Impact	Planned Fix
No token refresh	UX friction (30-min sessions)	Implement refresh tokens
Manual imports	Data staleness	Celery/background scheduler
Rate limit in-memory	No shared state across workers	Redis-backed counter
No CSP header	XSS mitigation incomplete	Add Content-Security-Policy

Compliance Checklist

Brief Requirement	Status	Location
GitHub repository	✓	github.com/NathS04/comp3011-cw1-api
API documentation PDF	✓	docs/API_DOCUMENTATION.pdf
Technical report PDF (≤ 5 pages)	✓	TECHNICAL_REPORT.pdf
Presentation slides	✓	docs/PRESENTATION_SLIDES.pptx
GenAI logs appendix	✓	docs/GENAI_EXPORT_LOGS.pdf
README.md	✓	Root directory
Deployed API URL	✓	comp3011-cw1-api.onrender.com

References

- [1] UK Government, “Open Government Licence v3.0,” 2014. Available: <https://www.nationalarchives.gov.uk/doc/open-government-licence/version/3/>
 - [2] Auth0, “JSON Web Tokens Best Practices,” 2023. Available: <https://auth0.com/blog/a-look-at-the-latest-draft-for-jwt-bcp/>
 - [3] Alembic, “Alembic Documentation,” 2024. Available: <https://alembic.sqlalchemy.org/en/latest/>
 - [4] FastAPI, “FastAPI Documentation,” 2024. Available: <https://fastapi.tiangolo.com/>
 - [5] SQLAlchemy, “SQLAlchemy 2.0 Documentation,” 2024. Available: <https://docs.sqlalchemy.org/en/20/>
 - [6] Render, “Render Documentation,” 2024. Available: <https://render.com/docs>
-

Report for COMP3011 CW1, University of Leeds