

# Technical Report

## COMP3011 Technical Report: EventHub – Event & RSVP API

---

**Module:** COMP3011 – Web Services and Web Data

**Coursework:** CW1 – Individual Web Services API Development

**Student:** Nathaniel Sebastian (sc232ns)

**Date:** 4th February 2026

**GitHub Repository:** [github.com/NathS04/comp3011-cw1-api](https://github.com/NathS04/comp3011-cw1-api)

**Live API:** [comp3011-cw1-api.onrender.com](https://comp3011-cw1-api.onrender.com)

**API Documentation (PDF):** See `docs/API_DOCUMENTATION.pdf`

---

## Abstract

EventHub is a RESTful API for managing event registrations, designed for student societies and community organisations. Beyond standard CRUD functionality, this project implements **novel data integration** through an idempotent dataset import pipeline with full provenance tracking, and provides **analytics endpoints** for seasonality analysis, trending detection, and personalised recommendations. The system demonstrates clean architectural separation, comprehensive test coverage (25 tests), and thoughtful use of Generative AI as a development partner. This report documents the design decisions, implementation challenges, and critical reflections on AI-assisted development.

---

## 1. Problem & Scope

### 1.1 Problem Domain

Event registration is a universal need—from university societies to corporate conferences. Organisers need to create events, track capacity, and understand attendance patterns. Attendees want to discover relevant events and RSVP easily.

This project builds a backend API for this domain, prioritising:

- Clean data modelling with proper constraints
- Secure authentication
- Extensibility for analytics and external data

## 1.2 Functional Requirements

Requirement	Implementation
Event CRUD	Create, read, update, delete events with validation
Attendee management	Register attendees with unique emails
RSVP tracking	Link attendees to events with status (going/maybe/not_going)
Capacity statistics	Calculate remaining spots and RSVP breakdown
Authentication	JWT-based auth for write operations
<b>Novel data integration</b>	Import external datasets with provenance
<b>Analytics</b>	Seasonality, trending, and recommendations

## 1.3 Non-Goals (Explicit Exclusions)

To maintain scope, the following are **not** implemented:

- Real-time notifications or WebSocket support
- Payment processing or ticketing
- Mobile app or frontend (API-only)
- Machine learning models (recommendations use deterministic scoring)
- Role-based access control (all authenticated users have equal permissions)


# 2. Architecture

## 2.1 Technology Stack

Component	Technology	Justification
<b>Framework</b>	FastAPI	Modern async support, automatic OpenAPI docs, excellent Pydantic integration
<b>Database</b>	SQLite (dev)	Zero-config, file-based. Prod-ready with PostgreSQL via DATABASE_URL
<b>ORM</b>	SQLAlchemy 2.x	Industry-standard, supports relationships and migrations
<b>Migrations</b>	Alembic	Version-controlled schema changes
<b>Auth</b>	JWT (python-jose)	Stateless tokens, fits REST principles

Component	Technology	Justification
Testing	pytest + TestClient	Isolated in-memory DB with StaticPool

## 2.2 Layered Architecture





**Key Principle:** Routes are thin handlers. All business logic lives in `crud.py`, making it testable and reusable.

---

## 3. Data Model

### 3.1 Entity-Relationship Diagram





### 3.2 Novel Tables: Data Provenance

To support external dataset integration, I added two tables:

Table	Purpose
data_sources	Tracks where external data came from (name, URL, license)
import_runs	Logs each import execution (rows processed, errors, status)

Events can optionally link to a `data_source` via `source_id`, enabling provenance queries.

## 4. Key Design Decisions

## Decision 1: RSVP as a First-Class Entity

**Choice:** Separate rsysps table with its own attributes (status, created\_at).

**Alternative considered:** Store RSVPs as a list field embedded in Event.

**Trade-off:** The separate table adds a JOIN for queries but enables:  
- Tracking when someone RSVP'd  
- Querying RSVPs independently (e.g., “all events user X is attending”)  
- Database-level uniqueness constraint

*I explored the embedded approach with AI assistance but found it created data integrity risks and made analytics harder.*

## Decision 2: JWT vs Server-Side Sessions

**Choice:** Stateless JWT tokens.

**Alternative considered:** Server-side session storage with Redis.

**Trade-off:** JWT is simpler for a pure API (no session store needed) but:  
- Tokens can't be revoked until expiry  
- Token size grows with claims

For this coursework scope, JWT simplicity outweighs revocation limitations.

## Decision 3: SQLite with PostgreSQL Path

**Choice:** SQLite for development, designed for PostgreSQL in production.

**Alternative considered:** PostgreSQL from the start.

**Trade-off:** SQLite enables zero-configuration local development. The code uses standard SQLAlchemy patterns that work with both databases. Alembic migrations use `batch_alter_table` for SQLite compatibility.

---

# 5. Novel Data Integration

## 5.1 Dataset Source

The system supports importing external event data. For demonstration, I created a sample CSV representing events from a fictional “Leeds Public Events API”:

Field	Example
id	EXT_001
title	Advanced AI Workshop
location	Leeds Tech Hub
start_time	2026-05-10T10:00:00
capacity	50

**License:** Sample data is CC-BY-4.0 compliant.

## 5.2 Import Pipeline Design

The `scripts/import_dataset.py` script implements:

1. **Source registration:** Creates or retrieves DataSource record
2. **Run logging:** Creates ImportRun with status “running”
3. **Idempotent import:** Uses source\_record\_id to detect existing records
4. **Error collection:** Continues on row errors, logs to errors\_json
5. **Summary output:** Reports rows read, inserted, updated, errors

```
# Idempotency check
existing = db.query(Event).filter(
    Event.source_id == source.id,
    Event.source_record_id == record_id
).first()

if existing:
    # Update existing record
    rows_updated += 1
else:
    # Insert new record
    rows_inserted += 1
```

## 5.3 Verification Output

---

```
Created new data source: Leeds Public Events API
Started Import Run ID: 1
Import Finished. Status: success
Read: 3, Inserted: 3, Updated: 0, Errors: 0
```

---

# 6. Analytics & Recommendations

## 6.1 Seasonality Endpoint

**Purpose:** Aggregate events by month to identify peak periods.

**Endpoint:** GET /analytics/events/seasonality

**Response:** Monthly counts with top categories (placeholder for future enhancement).

## 6.2 Trending Events

**Purpose:** Surface events gaining momentum.

**Endpoint:** GET /analytics/events/trending?window\_days=30

**Scoring Formula:**

$\text{trending\_score} = (\text{recent\_rsvps} \times 1.5) + (\text{total\_rsvps} \times 0.5)$

The 1.5x multiplier for recent RSVPs rewards events with growing interest.

## 6.3 Personalised Recommendations

**Purpose:** Suggest relevant events based on user history.

**Endpoint:** GET /events/recommendations

**Algorithm:** 1. Find attendee record for authenticated user (by email) 2. Extract locations from past RSVPs 3. Recommend future events at those locations 4. Cold start: Return top upcoming events if no history

This is deterministic scoring, not machine learning—appropriate for coursework scope while demonstrating the concept.

---

## 7. Testing Strategy

### 7.1 Approach

- **Isolated database:** In-memory SQLite with `StaticPool`
- **Fresh state:** Tables recreated per test function
- **Both paths tested:** Success cases and error handling (401, 404, 409, 422)

### 7.2 Coverage Summary

Test File	Tests	What's Covered
<code>test_auth.py</code>	6	Register, login, wrong password, protected routes
<code>test_events.py</code>	5	CRUD operations, pagination, filtering
<code>test_attendees.py</code>	4	Create, get, list events for attendee
<code>test_rsvps.py</code>	4	Create, list, delete, duplicate rejection
<code>test_analytics.py</code>	4	Seasonality, trending, recommendations
<code>test_import.py</code>	1	Data provenance model constraints
<code>test_health.py</code>	1	Health endpoint

**Total: 25 tests, all passing.**

---

## 8. Deployment

### 8.1 Platform

Deployed on Render.com with GitHub integration for automatic deploys.

## 8.2 Environment Variables

Variable	Purpose
DATABASE_URL	PostgreSQL connection (Render-provided)
SECRET_KEY	JWT signing key
ENVIRONMENT	production

## 8.3 Verification

- Health check: GET /health returns {"ok": true}
- Swagger UI: /docs accessible
- Sample requests via curl verified

# 9. Generative AI Usage Declaration

## 9.1 Tools Used

Tool	Purpose
Google Gemini (Antigravity)	Primary assistant—architecture, code generation, debugging
ChatGPT (GPT-4)	Secondary—marking feedback interpretation

## 9.2 How AI Was Used

My workflow integrated AI throughout:

1. **Architecture planning:** Described requirements, AI suggested layered structure
2. **Scaffolding:** Generated initial models, schemas, CRUD functions
3. **Debugging:** Pasted tracebacks, received explanations and fixes
4. **Documentation:** Drafted sections, edited for accuracy and my voice

## 9.3 Exploring Alternatives (Outstanding-Level Usage)

Rather than just accepting AI suggestions, I used it to explore design alternatives:

### Example 1: RSVP Storage

**Me:** Should I store RSVPs as a list in the Event model or as a separate table?

**AI:** A separate table enables uniqueness constraints, independent queries, and timestamps. Embedded lists are simpler but create data integrity risks.

**My decision:** Separate table—alignment with relational design principles.

### Example 2: Auth Approach

**Me:** JWT vs session-based auth for a REST API?

**AI:** JWT is stateless and simpler for APIs. Sessions require storage but enable revocation.

**My decision:** JWT for simplicity, with documented limitation about revocation.

### Example 3: SQLite Compatibility

**AI initially suggested:** Standard `op.add_column` for migrations.

**Problem:** SQLite doesn't support `ALTER TABLE` for foreign keys.

**Solution found with AI:** Use Alembic's `batch_alter_table` which recreates the table.

## 9.4 Bugs from AI Suggestions

1. **bcrypt compatibility:** AI suggested bcrypt, but my environment had issues.  
Switched to pbkdf2\_sha256.
2. **Circular imports:** AI-generated auth code caused import cycles. Fixed with local imports.
3. **Test fixtures:** Initial fixture didn't isolate state properly. Redesigned with per-test table recreation.

## 9.5 Critical Reflection

AI accelerated development significantly—especially for boilerplate and debugging. However:

- I always ran generated code before committing
- I caught security issues (hardcoded secrets) in review
- I ensured I understood every line before accepting it

AI was a partner, not an author. The architectural decisions and critical evaluation were mine.

---

# 10. Limitations & Future Work

## Current Limitations

- Single-tenant: All authenticated users can modify all events
- No email verification for registration
- Token expiry (30 min) with no refresh mechanism
- SQLite concurrency limitations

## Future Roadmap

- PostgreSQL for production scalability
  - Role-based access (admin vs. user)
  - Capacity enforcement (reject RSVPs when full)
  - Email notifications
  - Rate limiting
  - Event categories for richer analytics
- 

## Appendix A: GenAI Conversation Logs

*See attached docs/appendix\_genai\_logs.md for selected excerpts.*

---

**Word Count:** ~2,100 words (excluding diagrams and code)

*Report generated for COMP3011 CW1 submission, University of Leeds*