

COMP3011 Technical Report: EventHub API

Module: COMP3011 – Web Services and Web Data

Student: Nathaniel Sebastian (sc232ns)

Date: 5th February 2026

Word Count: ~1,800 words

Submission Links

Resource	Link
GitHub Repository	github.com/NathS04/comp3011-cw1-api
Live API	comp3011-cw1-api.onrender.com

1. Compliance Checklist

Requirement	Status	Location
GitHub repository	✓	github.com/NathS04/comp3011-cw1-api
API documentation PDF	✓	docs/API_DOCUMENTATION.pdf
Technical report PDF	✓	TECHNICAL_REPORT.pdf
Presentation slides	✓	docs/PRESENTATION_SLIDES.pptx
GenAI logs	✓	docs/GENAI_EXPORT_LOGS.pdf
README.md	✓	Root directory
Deployed API	✓	comp3011-cw1-api.onrender.com
Novel data integration	✓	Leeds TEN XML with SHA256 provenance
Authentication	✓	JWT + PBKDF2
Test suite	✓	41 tests passing

Marker Evidence Cross-Reference

Requirement	Evidence	File Path
Tests pass (41)	<code>pytest -q output</code>	VERIFICATION_INSTRUCTIONS.md

Requirement	Evidence	File Path
Live deployment	/health → environment: prod	comp3011-cw1-api.onrender.com/health
API documentation	Full endpoint reference	docs/API_DOCUMENTATION.pdf
Security headers	7 headers on every response	app/core/middleware.py
Rate limiting	120/min global, 10/min login	app/core/rate_limit.py
Novel dataset	Leeds TEN XML, SHA256 hashed	scripts/import_dataset.py , docs/DATASET_SOURCE.md
GenAI logs	Tools, failures, dates	docs/GENAI_EXPORT_LOGS.pdf
RBAC	Admin-only /admin/* with 403	app/api/admin.py , app/core/auth.py
ETag/304	RFC 7232 compliance	app/core/middleware.py
Presentation	Slide outline + viva prep	docs/PRESENTATION_OUTLINE.md

2. Reproducibility

```
git clone https://github.com/NathS04/comp3011-cw1-api.git && cd comp3011-cw1-api
python3 -m venv venv && source venv/bin/activate
pip install -r requirements.txt
export DATABASE_URL="sqlite:///./app.db" SECRET_KEY="dev-secret"
alembic upgrade head
pytest -q                      # Expected: 41 passed
uvicorn app.main:app --reload
```

3. Dataset Provenance

Attribute	Value
Source	Leeds Temporary Event Notices [1]
Provider	Leeds City Council (Data Mill North)
Licence	Open Government Licence v3.0
Format	XML
Fields Mapped	Reference_Number → ID, Premises_Name → Title

Attribute	Value
Provenance	SHA256 hash stored in <code>ImportRun</code> table

4. Architecture

```
Client → Middleware (Rate Limit, Headers, ETag) → FastAPI → Auth → Routes → CRUD → SQLAlchemy → DB
```

Design Principles: - Layered architecture: thin routes, fat CRUD - Middleware handles cross-cutting concerns (rate limiting, security headers, ETag) - Dual database support via Alembic (SQLite dev, PostgreSQL prod)

5. Outstanding Evidence

5.1 Test Coverage

Category	Count	Coverage
Auth	6	Register, login, validation
Events CRUD	5	Create, read, update, delete, pagination
RSVPs	4	Create, duplicate rejection
Analytics	4	Seasonality, trending, recommendations
Admin/Import	3	Idempotency, provenance
RBAC	2	403 for non-admin
Attendees	4	CRUD, uniqueness
Health	1	Returns metadata
ETag	3	Generation, 304, mismatch
Security Headers	3	On 200, 404, 403
Error Handling	2	Sanitization, request_id
Total	41	

Test isolation: In-memory SQLite with `StaticPool`; tables created/dropped per test. Runtime: <1.5s.

5.2 Security Header Policy

All responses include:

Header	Value	OWASP Reference
X-Request-ID	UUID v4	Tracing/debugging
X-Content-Type-Options	nosniff	MIME sniffing prevention [2]
X-Frame-Options	DENY	Clickjacking protection [2]
Referrer-Policy	no-referrer	Privacy [2]
Permissions-Policy	geolocation=(), microphone=(), camera=()	Feature restriction [2]
Cross-Origin-Resource-Policy	same-site	Side-channel protection [2]
Cache-Control	no-store (default) / no-cache (ETag)	Sensitive data protection

5.3 Rate Limiting

Scope	Limit	Rationale
Global	120 req/min	Prevents abuse without impacting normal use
/auth/login	10 req/min	Credential stuffing mitigation [3]

429 response includes `request_id` for correlation: `{"detail": "Too Many Requests", "request_id": "<uuid>"}`

5.4 ETag/304 Standard Compliance

Implementation follows RFC 7232 [4]:

1. GET `/events` and `/events/{id}` compute SHA256 of response body
2. Response includes `ETag: "<hash>"`
3. Client sends `If-None-Match: "<hash>"`
4. Server returns 304 Not Modified with empty body if match
5. 304 still includes `ETag`, `X-Request-ID`, and security headers

5.5 Provenance Hashing

Each import run stores: - SHA256 hash of source file - Timestamp, duration, row counts - Enables idempotency (duplicate hash = skip)

6. Design Trade-offs

Decision	Alternative	Trade-off	Justification
In-memory rate limiting	Redis	Single-process only	Acceptable for free-tier Render; documented as limitation
ETag via body hash	DB <code>updated_at</code> column	Recalculates per request	Simpler implementation; no schema changes
JWT without refresh	Refresh token flow	30-min hard limit	Simpler for coursework; users re-login
SQLite/PostgreSQL dual	PostgreSQL-only	Dev overhead	Alembic abstracts dialect; faster local iteration
Sanitized 500 errors	Detailed error messages	Less debugging info	Security first; <code>request_id</code> enables log correlation

7. Analytics Endpoints

Endpoint	Computation	Use Case
<code>/analytics/events/seasonality</code>	COUNT(*) GROUP BY month	Identify peak event periods
<code>/analytics/events/trending</code>	$(\text{recent_rsvps} \times 1.5) + (\text{total_rsvps} \times 0.5)$	Surface popular events
<code>/events/recommendations</code>	Filter by user's past RSVP categories	Personalised discovery

Trade-off Rationale: Simpler algorithms chosen for coursework scope; documented as extensible via Redis caching and ML in production.

8. Security Implementation

Security Measure	Implementation	File
Password Hashing	PBKDF2-SHA256	<code>app/core/auth.py</code>
JWT Signing	HS256, 30-min expiry	<code>app/core/auth.py</code>
RBAC	<code>is_admin</code> flag, 403 on <code>/admin/*</code>	<code>app/core/auth.py</code> , <code>app/api/admin.py</code>
Rate Limiting	120/min global, 10/min login	<code>app/core/rate_limit.py</code>

Security Measure	Implementation	File
Request Tracing	X-Request-ID on all responses	app/core/middleware.py
Security Headers	6 headers on all responses	app/core/middleware.py
ETag Caching	SHA256 body hash, 304 support	app/core/middleware.py
Error Sanitization	Generic 500 message, no stack trace	app/core/middleware.py

9. Deployment

Platform: Render.com [5]

Database: Managed PostgreSQL

Config: render.yaml

Variables: DATABASE_URL , SECRET_KEY , ENVIRONMENT=prod , ALLOWED_ORIGINS , RATE_LIMIT_ENABLED

Cold Start Note: Free tier spins down after 15 min; first request may take 30–60s.

9.1 Deployment Verification

Live health check confirms production configuration:

```
curl https://comp3011-cw1-api.onrender.com/health
{
  "status": "online",
  "database": "ok",
  "version": "1.0.0",
  "environment": "prod",
  "commit": "<current HEAD>",
  "timestamp": "2026-02-09T..."
}
```

Environment detection logic (app/core/config.py):- If RENDER env var exists → always "prod" (cannot be overridden by .env) - load_dotenv() is skipped entirely on Render to prevent stale config leakage - Locally: respects ENVIRONMENT env var, defaults to "dev"

Render environment variables (via render.yaml + dashboard):

Variable	Value	Source
ENVIRONMENT	prod	render.yaml
DATABASE_URL	PostgreSQL connection string	Render managed DB
SECRET_KEY	Auto-generated	render.yaml
RATE_LIMIT_ENABLED	true	render.yaml

Variable	Value	Source
ALLOWED_ORIGINS	Restricted to Render + localhost	render.yaml

10. Version Control

Repository: github.com/NathS04/comp3011-cw1-api

Commit Discipline

Development followed incremental feature-branch-style commits on `main`, with meaningful conventional-commit messages (`feat:`, `fix:`, `docs:`, `test:`, `chore:`). Each commit represents a logically atomic change verified by the test suite.

```
51c27c1 fix(test): use fixture in header tests + hardened zip exclusions
dca5aa6 docs(claude): regenerate PDFs + add marker evidence checklist
23b7054 fix(test): use fixture for etag tests
f72a36b fix(config): allow rate limit true/1 and clean zip
fa1f4e5 docs(verify): bulletproof verification steps
6d2a33d docs(slides): align narrative + viva prep
acfb22c docs(genai): final audit-ready logs
5d31b97 docs(report): finalize outstanding evidence + tradeoffs
38168e5 docs(api): finalize errors/headers/etag/flows
f9d7fe6 docs(readme): sync final features + commands
cfaca6c test(final): ensure >=39 tests and green
ef8a935 chore(devtools): add ruff/mypy/bandit/pip-audit
342473d fix(errors): sanitize all 500s, no leakage
37d5574 feat(etag): ETag + 304 for list + detail
f52d165 fix(rate): 429 includes request_id + headers
```

Observations: - 45+ commits total, demonstrating iterative development - Clear separation of concern: feature commits, bug fixes, documentation, and testing - No monolithic dump commits — each change is reviewable in isolation

11. GenAI Usage

Tools: Google Gemini (Antigravity), Claude, ChatGPT

Creative Applications: 1. Compared RSVP storage approaches (embedded vs relational) 2. Evaluated rate limiting options (in-memory vs Redis) [3] 3. Researched ETag/If-None-Match RFC 7232 compliance [4] 4. Compared middleware ordering alternatives for header injection

Failures Caught: - Missing `requests` dependency → `ModuleNotFoundError` - Placeholder test with `pass` → False coverage - Deprecated `Query(regex=...)` → Warning - Headers not applied to 429 responses → Fixed middleware flow

Full logs: [docs/GENAI_EXPORT_LOGS.pdf](#)

12. Limitations & Future Work

Limitation	Impact	Improvement
No token refresh	Users re-login after 30 min	Implement refresh tokens
In-memory rate limiting	Resets on restart; single-worker only	Redis-backed for horizontal scaling
No CSP header	Reduced browser-side protection	Add Content-Security-Policy
Basic recommendations	Location-based scoring only	Collaborative filtering with user embeddings
Analytics via Python loops	$O(n)$ over all events	SQL aggregates (GROUP BY, window functions)

Rate limiter justification: Render free tier runs a single worker process, so in-memory state is not lost to multi-instance divergence. The dictionary-based approach is appropriate for the deployment target. Redis migration is documented as the production upgrade path.

References

- [1] Leeds City Council, "Temporary Event Notices," Data Mill North, <https://datamillnorth.org/dataset/temporary-event-notices> (Open Government Licence v3.0)
 - [2] OWASP, "HTTP Security Response Headers," https://cheatsheetseries.owasp.org/cheatsheets/HTTP_Headers_Cheat_Sheet.html
 - [3] OWASP, "Rate Limiting," https://cheatsheetseries.owasp.org/cheatsheets/Denial_of_Service_Cheat_Sheet.html
 - [4] IETF, "RFC 7232: HTTP/1.1 Conditional Requests," <https://tools.ietf.org/html/rfc7232>
 - [5] Render, "Web Services Documentation," <https://render.com/docs/web-services>
-