

COMP3011 Technical Report: EventHub API

Module: COMP3011 – Web Services and Web Data

Student: Nathaniel Sebastian (sc232ns)

Date: 5th February 2026

GitHub: github.com/NathS04/comp3011-cw1-api

Live API: comp3011-cw1-api.onrender.com

1. Reproducibility

Quickstart (Fresh Clone): 1. `python3 -m venv .venv && source .venv/bin/activate` 2. `pip install -r requirements.txt` 3. `alembic upgrade head` 4. `python -c "import app.main"` (Verifies imports) 5. `pytest -q` (Runs 31 tests) 6. `uvicorn app.main:app --reload`

Expected Results: all 31 tests passed. API available at <http://127.0.0.1:8000/docs>.

Note: `scripts/import_dataset.py` requires an internet connection to fetch the XML feed.

2. Dataset Provenance

Attribute	Details
Source Name	Leeds Temporary Event Notices (TENs)
Provider	Leeds City Council (Data Mill North)
Format	XML (Live Feed)
URL	https://opendata.leeds.gov.uk/downloads/Licences/temp-event-notice/temp-event-notice.xml
Last Updated	Daily (as per portal metadata)
Retrieval Date	5th February 2026
Fields Used	<code>Reference_Number</code> (ID), <code>Premises_Name</code> (Title), <code>Activities</code> (Desc), <code>Event_Start_Date</code> , <code>Event_End_Date</code>
Limitations	No geolocation coordinates; Categories are unstructured text; Times often “00:00”

Why this dataset? It represents a “messy” real-world source requiring non-trivial XML parsing, date normalization (DD/MM/YYYY -> ISO8601), and robust error handling (skipping malformed records).

3. Architecture

The system follows a layered architecture to ensure separation of concerns and testability:

1. **Router Layer (app/api)**: Handles HTTP requests/responses, auth checks, and schema validation.
2. **Service/CRUD Layer (app/crud.py)**: Contains pure business logic. Decoupled from HTTP.
3. **Data Layer (app/models.py)**: SQLAlchemy ORM models mapping to database tables.
4. **Database**: SQLite (Development) / PostgreSQL (Production).

Diagram: Client (React/Curl) <-> FastAPI (Pydantic) <-> Logic (SQLAlchemy) <-> PostgreSQL

4. Engineering Trade-offs & Design Decisions

SQLite vs PostgreSQL

- **Decision:** Use SQLite locally, PostgreSQL on Render.
- **Trade-off:** SQLite is zero-config but lacks advanced concurrent writing. Postgres complicates dev setup but ensures production reliability.
- **Mitigation:** alembic handles dialect differences (e.g., SQLite's lack of ALTER TABLE).

XML Parsing Strategy (DOM vs Streaming)

- **Decision:** Used `xml.etree.ElementTree` (DOM-style).
- **Trade-off:** Loads entire file into memory. Fast for 500 records (<1MB), but risky for gigabyte-scale files.
- **Why acceptable:** Leeds TEN dataset is small (~50KB). If it grows >100MB, would switch to `iterparse`.

Auth Implementation

- **Decision:** Stateless JWT (HS256) with 30-minute expiry.
- **Trade-off:** Simple scaling (no Redis session store needed) vs inability to revoke tokens immediately.
- **Mitigation:** Documented short expiry as security control.

Recommendation Algorithm

- **Decision:** Simple location-matching heuristic.
- **Trade-off:** Computational speed ($O(1)$ lookup) vs “smartness” (no collaborative filtering).
- **Why acceptable:** Sufficient for cold-start demo; runs in <10ms.

5. Dataset Ingestion Pipeline

Script: scripts/import_dataset.py

1. **Fetch:** Downloads XML from Leeds Open Data.
2. **Verify:** Computes SHA256 hash of raw content for provenance.
3. **Parse:** Extract Temporary_Event_Notify nodes.
4. **Upsert:** Uses source_record_id (Reference No) to identify duplicates. Updates existing records instead of creating duplicates.
5. **Log:** Records usage stats to import_runs table (duration, rows inserted/updated).

Robustness: Continue-on-error behavior (skips single bad rows, logs error, proceeds).

6. Evaluation (Metrics)

Metric	Result	Environment
Import Time	~2.1 seconds	Local M1 Mac (WiFi)
Import Throughput	~240 records/sec	Local M1 Mac
API Latency (Read)	8ms (avg)	Local SQLite
API Latency (Write)	12ms (avg)	Local SQLite
Test Coverage	31 Tests	100% Pass Rate

Test Distribution: * **Auth:** 5 tests (Register, Login, Bad Token) * **CRUD:** 15 tests (Events, RSVPs, Attendees) * **Analytics:** 6 tests (Seasonality, Trending, Recs) * **Import/Admin:** 5 tests (XML parsing, Idempotency)

7. GenAI Usage (Critical Reflection)

Tools & Role

- **Gemini 3 Pro:** “Junior Developer” role (Drafting boilerplate, writing tests).
- **Claude Opus:** “Senior Reviewer” role (Refactoring, documentation polish).

Successes

- Gemini successfully suggested the `StaticPool` pattern for in-memory SQLite testing, solving a “thread check” error.
- Claude generated the 9-section report structure which improved document clarity.

Failures & Corrections (Manual Intervention)

1. **Bcrypt Compatibility:** AI suggested an argon2 hasing library that failed to compile on Render. I manually reverted to `passlib[bcrypt]` for stability.
2. **Duplicate Imports:** Reference code generation created circular imports between `models.py` and `schemas.py`. I manually broke the cycle using string forward references.
3. **Missing Dependency:** AI wrote a script using `requests` but forgot to add it to `requirements.txt`. I caught this via `ModuleNotFoundError` and fixed the build process.

Full logs available in `docs/GENAI_EXPORT_LOGS.pdf`

8. Limitations & Future Work

1. **Single Tenant Admin:** No role separation (all users are basic). *Plan: Add `is_admin` column.*
2. **Rate Limiting:** API vulnerable to DoS. *Plan: Add `slowapi` or Nginx limiting.*
3. **Data Freshness:** Imports are manual/triggered. *Plan: Cron job or Celery task.*
4. **Token Refresh:** Users forced to relogin every 30m. *Plan: Implement refresh tokens.*

Word Count: ~880 words

Report for COMP3011 CW1, University of Leeds